

# DSA210 TERM PROJECT

## INTRODUCTION

This project will be about investigating how data can be used to determine a patient's lung cancer status based on important medical signs. The goal is to find the most crucial elements that go into a diagnosis by looking for trends in patient data. Understanding which characteristics—such as demographics, mental difficulties, or chronic diseases—are most predictive and how a data-driven strategy might support early detection are the main goals of this study.

## Import necessary libraries and pull the data of the first dataset

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import chi2_contingency
```

In [2]:

```
file_path = "dataset.csv"
df = pd.read_csv(file_path)
df.head()
```

Out[2]:

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CHRONIC_DISEASE	FATIGUE	ALLI
0	M	65	1	1	1	2	2	1	
1	F	55	1	2	2	1	1	2	
2	F	78	2	2	1	1	1	2	
3	M	60	2	1	1	1	2	1	
4	F	80	1	1	2	1	1	2	

In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 16 columns):
#   Column              Non-Null Count  Dtype
---  -
0   GENDER              3000 non-null   object
1   AGE                 3000 non-null   int64
2   SMOKING             3000 non-null   int64
3   YELLOW_FINGERS      3000 non-null   int64
4   ANXIETY              3000 non-null   int64
```

```
5    PEER_PRESSURE      3000 non-null    int64
6    CHRONIC_DISEASE    3000 non-null    int64
7    FATIGUE            3000 non-null    int64
8    ALLERGY            3000 non-null    int64
9    WHEEZING           3000 non-null    int64
10   ALCOHOL_CONSUMING  3000 non-null    int64
11   COUGHING           3000 non-null    int64
12   SHORTNESS_OF_BREATH 3000 non-null    int64
13   SWALLOWING_DIFFICULTY 3000 non-null    int64
14   CHEST_PAIN         3000 non-null    int64
15   LUNG_CANCER        3000 non-null    object
```

dtypes: int64(14), object(2)

memory usage: 375.1+ KB

In [4]:

```
df.describe()
```

Out[4]:

	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CHRONIC_DISEASE	FATIGUE
count	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000
mean	55.169000	1.491000	1.514000	1.494000	1.499000	1.509667	1.489000
std	14.723746	0.500002	0.499887	0.500047	0.500082	0.499990	0.499000
min	30.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000
25%	42.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000
50%	55.000000	1.000000	2.000000	1.000000	1.000000	2.000000	1.000000
75%	68.000000	2.000000	2.000000	2.000000	2.000000	2.000000	2.000000
max	80.000000	2.000000	2.000000	2.000000	2.000000	2.000000	2.000000

Changing the number 1s and 2s into No's and Yes's, M to Male and F to Female for better visualization.

In [5]:


```
binary_columns = [
    "SMOKING", "YELLOW_FINGERS", "PEER_PRESSURE", "CHRONIC_DISEASE",
    "FATIGUE", "ALLERGY", "WHEEZING", "ALCOHOL_CONSUMING", "ANXIETY",
    "COUGHING", "SHORTNESS_OF_BREATH", "SWALLOWING_DIFFICULTY", "CHEST_PAIN"
]

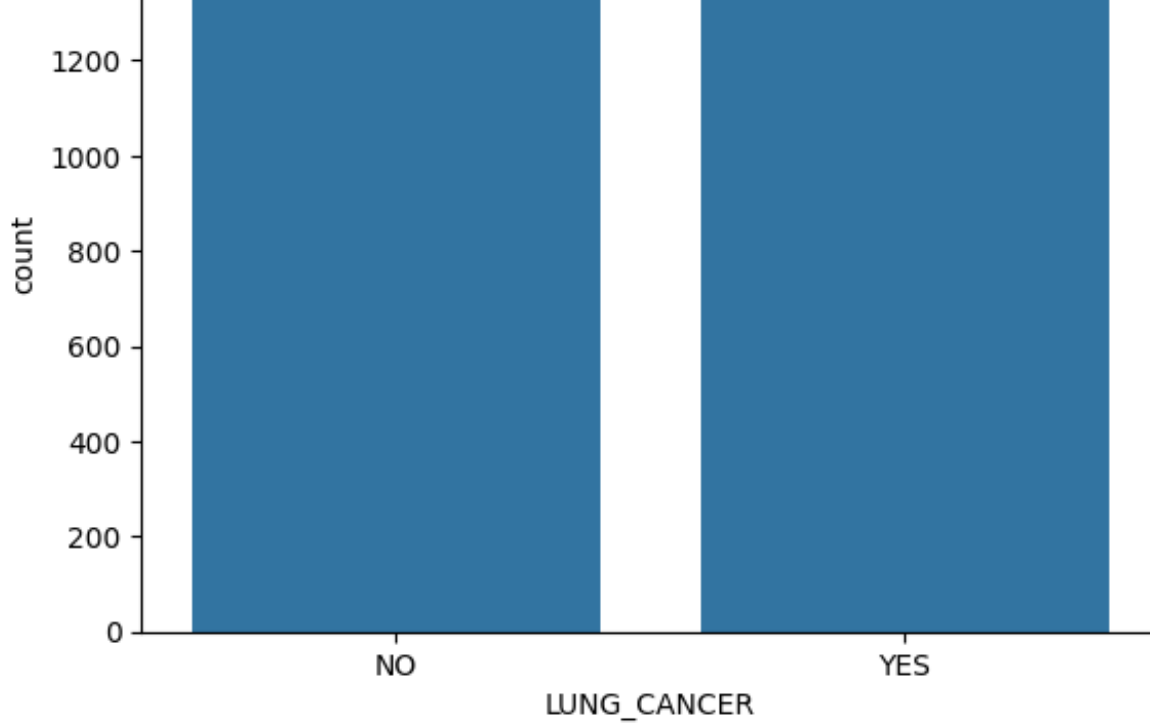
# Map 1 -> 'No', 2 -> 'Yes'
df[binary_columns] = df[binary_columns].replace({1: "No", 2: "Yes"})
df["GENDER"] = df["GENDER"].replace({"M": "Male", "F": "Female"})
```

Total distribution of cancer patients and healthy patients.

In [6]:

```
sns.countplot(x='LUNG_CANCER', data=df,)
plt.title('Target Distribution');
```






Plotting the most common assumptions for the causes of Lung Cancers.


In [7]:

```
plt.figure(figsize=(10, 5))
sns.countplot(x="SMOKING", hue="LUNG_CANCER", data=df)
plt.title("Smoking vs Lung Cancer")
plt.xlabel("Smoking")
plt.ylabel("Number of Patients")
plt.show()

plt.figure(figsize=(10, 5))
sns.countplot(x="CHRONIC_DISEASE", hue="LUNG_CANCER", data=df)
plt.title("Chronic Diseases vs Lung Cancer")
plt.xlabel("Chronic Disease")
plt.ylabel("Number of Patients")
plt.show()
```



Chronic Diseases vs Lung Cancer



### Age distributions of patients for Lung Cancer and Smoking

In [8]:


```
plt.figure(figsize=(10, 5))
sns.histplot(data=df, x="AGE", hue="LUNG_CANCER", kde=True, bins=50, palette="viridis")
plt.title("Age Distribution by Lung Cancer Diagnosis")
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.show()


plt.figure(figsize=(10, 5))
sns.histplot(data=df, x="AGE", hue="SMOKING", kde=True, bins=50, palette="coolwarm")
plt.title("Age Distribution by Smoking")
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.show()

mean_age = df['AGE'].mean()
median_age = df['AGE'].median()
mode_age = df['AGE'].mode()[0]

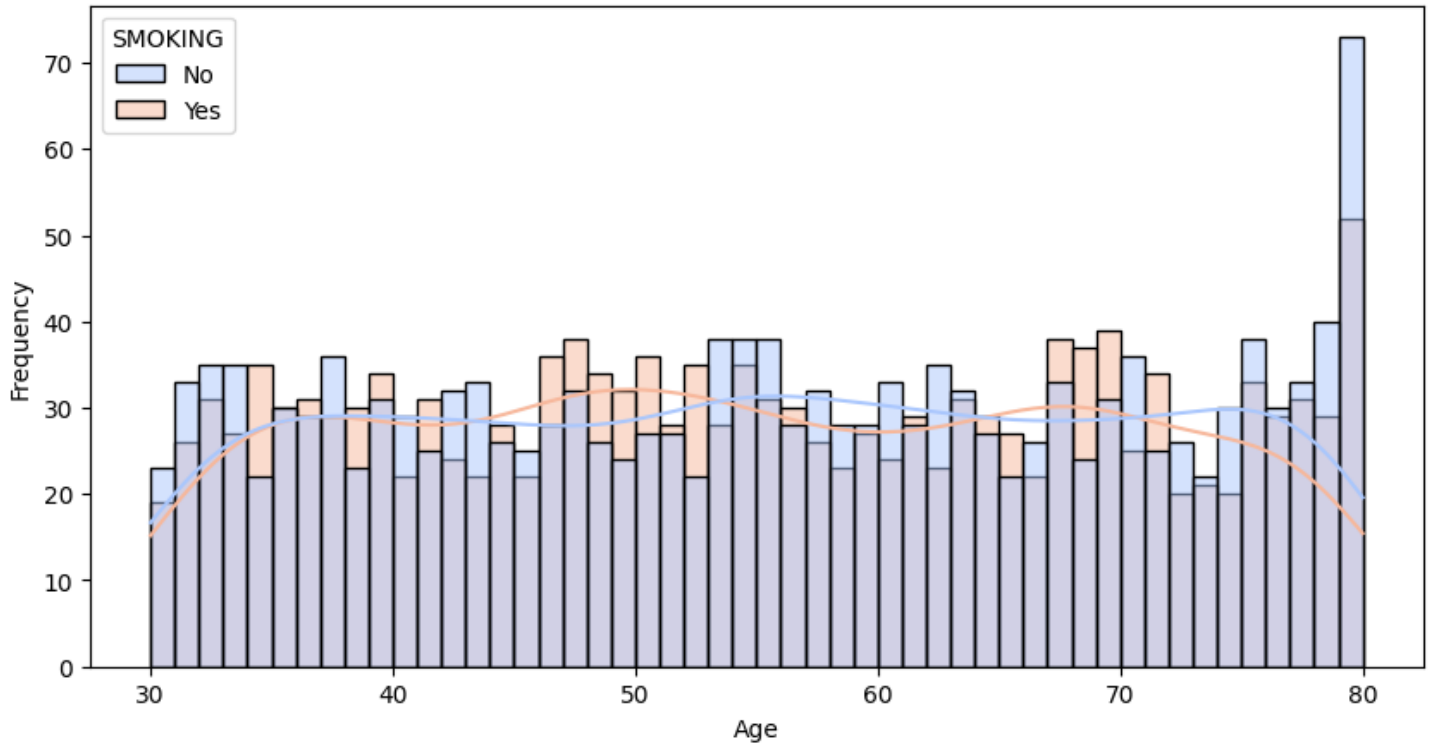
print(f"Mean Age: {mean_age}")
print(f"Median Age: {median_age}")
print(f"Mode Age: {mode_age}")
```

Age Distribution by Lung Cancer Diagnosis





Age Distribution by Smoking



Mean Age: 55.169  
Median Age: 55.0  
Mode Age: 54

## Male/Female Ratio

In [9]:

```
df["GENDER"] = df["GENDER"].replace({"M": "Male", "F": "Female"})


gender_counts = df["GENDER"].value_counts()

colors = ["#dcbef", "#e7e0ee"]
explode = (0.05, 0)

plt.figure(figsize=(6, 6))
plt.pie(
    gender_counts,
    labels=gender_counts.index,
    autopct="%1.1f%%",
    startangle=90,
    colors=colors,
    explode=explode,
    shadow=True
)
plt.title("Gender Distribution")
plt.axis("equal")
plt.show()
```

Gender Distribution

## Gender Distribution



## Some Columns of Possible Causes vs Age Using Violin Graphs

In [10]:

```
fig, axes = plt.subplots(nrows=2, ncols=3, figsize=(15, 10))

sns.violinplot(hue="LUNG_CANCER", y="AGE", data=df, inner="quartile", ax=axes[0][0],
palette="Reds", legend=False)
axes[0][0].set_title("Lung Cancer vs Age")

sns.violinplot(hue="SMOKING", y="AGE", data=df, inner="quartile", ax=axes[0][1], pal
ette="Blues", legend=False)
axes[0][1].set_title("Smoking vs Age")


sns.violinplot(hue="ALCOHOL_CONSUMING", y="AGE", data=df, inner="quartile", ax=axes[
0][2], palette="Greens", legend=False)
axes[0][2].set_title("Alcohol Consumption vs Age")

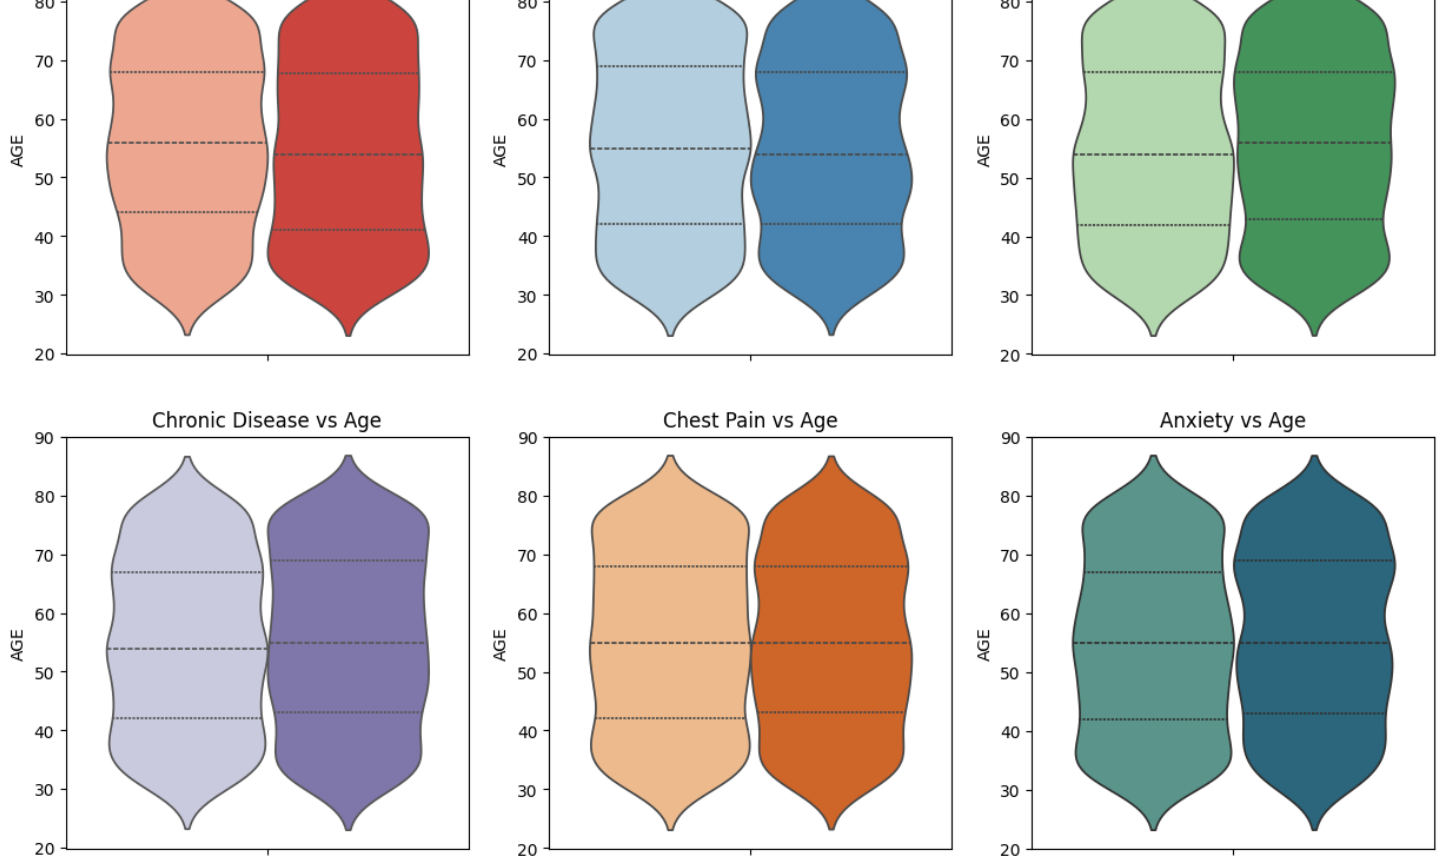
sns.violinplot(hue="CHRONIC_DISEASE", y="AGE", data=df, inner="quartile", ax=axes[1]
[0], palette="Purples", legend=False)
axes[1][0].set_title("Chronic Disease vs Age")

sns.violinplot(hue="CHEST_PAIN", y="AGE", data=df, inner="quartile", ax=axes[1][1],
palette="Oranges", legend=False)
axes[1][1].set_title("Chest Pain vs Age")

sns.violinplot(hue="ANXIETY", y="AGE", data=df, inner="quartile", ax=axes[1][2], pal
ette="crest", legend=False)
axes[1][2].set_title("Anxiety vs Age")

plt.show()
```





## All Columns vs Lung Cancer

In [11]:

```
features = [
    "SMOKING", "YELLOW_FINGERS", "ANXIETY", "PEER_PRESSURE", "CHRONIC_DISEASE",
    "FATIGUE", "ALLERGY", "WHEEZING", "ALCOHOL_CONSUMING", "COUGHING",
    "SHORTNESS_OF_BREATH", "SWALLOWING_DIFFICULTY", "CHEST_PAIN", "GENDER", "AGE"
]

fig, axes = plt.subplots(3, 5, figsize=(17.5, 10))
axes = axes.flatten()


for i, feature in enumerate(features):
    if feature == "AGE":
        axes[i].axis('off')
        continue
    crosstab = pd.crosstab(df[feature], df["LUNG_CANCER"], normalize='index')
    crosstab.plot(kind='bar', stacked=True, ax=axes[i], color=["#F44336", "#4CAF50"])

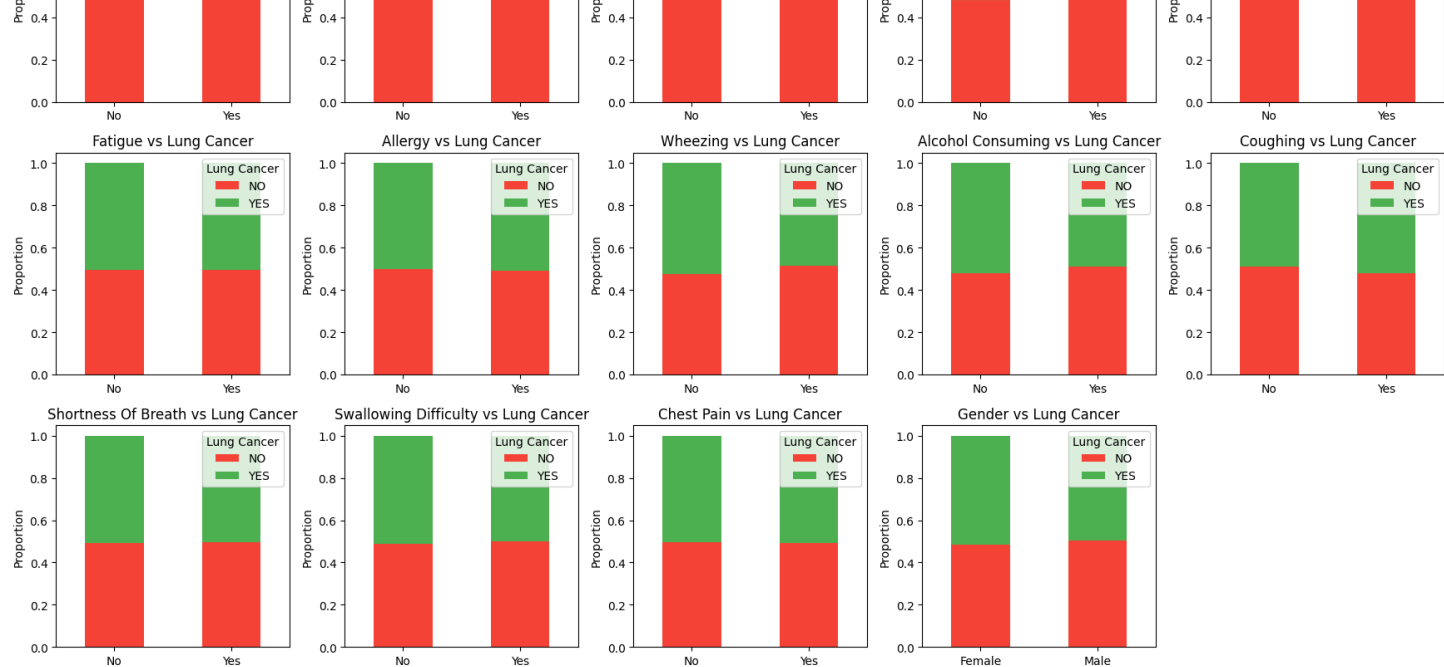
    axes[i].set_title(f"{feature.replace('_', ' ').title()} vs Lung Cancer")
    axes[i].set_ylabel("Proportion")
    axes[i].set_xlabel("")
    axes[i].legend(title="Lung Cancer", loc='upper right')

    axes[i].set_xticklabels(axes[i].get_xticklabels(), rotation=0)

for j in range(i+1, len(axes)):
    axes[j].axis('off')

plt.tight_layout()
plt.show()
```





## Correlation Matrix of Lung Cancer Dataset 1

In [12]:

```
df_encoded = df.copy()
yes_no_map = {"Yes": 1, "No": 0}
gender_map = {"Male": 1, "Female": 0}

yes_no_columns = [
    "SMOKING", "YELLOW_FINGERS", "ANXIETY", "PEER_PRESSURE",
    "CHRONIC_DISEASE", "FATIGUE", "ALLERGY", "WHEEZING", "ALCOHOL_CONSUMING",
    "COUGHING", "SHORTNESS_OF_BREATH", "SWALLOWING_DIFFICULTY", "CHEST_PAIN", "LUNG_
    CANCER"
]

df_encoded[yes_no_columns] = df_encoded[yes_no_columns].map(yes_no_map.get)
df_encoded["GENDER"] = df_encoded["GENDER"].map(gender_map)

df_numeric = df_encoded.select_dtypes(include=["number"])

corr_matrix = df_numeric.corr()


plt.figure(figsize=(13, 10))
sns.heatmap(
    corr_matrix,
    annot=True,
    fmt=".2f",
    cmap="RdYlBu",
    center=0,
    linewidths=0.5,
    square=True,
    cbar_kws={"shrink": 0.75}
)

plt.title("Correlation Matrix of Lung Cancer Dataset", fontsize=16)
plt.xticks(rotation=45)
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()
```

Correlation Matrix of Lung Cancer Dataset

GENDER	1.00	0.01	-0.03	0.01	-0.02	0.01	0.01	0.00	0.01	-0.01	-0.01	-0.01	0.01	-0.02	-0.00
AGE	0.01	1.00	-0.02	0.02	0.03	0.00	-0.03	-0.00	0.02	-0.02	-0.02	-0.02	-0.02	-0.00	-0.01





## Pulling the data from the second dataset and visualizing it


In [13]:

```
file_path2 = "dataset2.csv"
df2 = pd.read_csv(file_path2)
df2.head()
```

Out[13]:

index	Patient Id	Age	Gender	Air Pollution	Alcohol use	Dust Allergy	OccuPational Hazards	Genetic Risk	chronic Lung Disease	...	Fatigue	Weight Loss
0	0	P1	33	1	2	4	5	4	3	2	3	4
1	1	P10	17	1	3	1	5	3	4	2	1	3
2	2	P100	35	1	4	5	6	5	5	4	8	7
3	3	P1000	37	1	7	7	7	7	6	7	4	2
4	4	P101	46	1	6	8	7	7	7	6	3	2

5 rows x 26 columns



Removing the index and Patient Id columns

In [14]:

```
df2 = df2.iloc[:, 2:]
```

In [15]:

```
df2.head()
```

Out[15]:

	Age	Gender	Air Pollution	Alcohol use	Dust Allergy	OccuPational Hazards	Genetic Risk	chronic Lung Disease	Balanced Diet	Obesity	...	Fatigue	Weigh Loss
0	33	1	2	4	5	4	3	2	2	4	...	3	
1	17	1	3	1	5	3	4	2	2	2	...	1	
2	35	1	4	5	6	5	5	4	6	7	...	8	
3	37	1	7	7	7	7	6	7	7	7	...	4	
4	46	1	6	8	7	7	7	6	7	7	...	3	

5 rows x 24 columns



In [16]:

```
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1000 non-null   int64
1   Gender                                1000 non-null   int64
2   Air Pollution                         1000 non-null   int64
3   Alcohol use                           1000 non-null   int64
4   Dust Allergy                          1000 non-null   int64
5   OccuPational Hazards                  1000 non-null   int64
6   Genetic Risk                          1000 non-null   int64
7   chronic Lung Disease                  1000 non-null   int64
8   Balanced Diet                         1000 non-null   int64
9   Obesity                               1000 non-null   int64
10  Smoking                               1000 non-null   int64
11  Passive Smoker                        1000 non-null   int64
12  Chest Pain                            1000 non-null   int64
13  Coughing of Blood                     1000 non-null   int64
14  Fatigue                               1000 non-null   int64
15  Weight Loss                           1000 non-null   int64
16  Shortness of Breath                   1000 non-null   int64
17  Wheezing                              1000 non-null   int64
18  Swallowing Difficulty                 1000 non-null   int64
19  Clubbing of Finger Nails              1000 non-null   int64
20  Frequent Cold                         1000 non-null   int64
21  Dry Cough                             1000 non-null   int64
22  Snoring                               1000 non-null   int64
23  Level                                 1000 non-null   object
dtypes: int64(23), object(1)
memory usage: 187.6+ KB
```

In [17]:

```
df2.iloc[:, 0:11].describe()
```

Out[17]:

	Age	Gender	Air Pollution	Alcohol use	Dust Allergy	OccuPational Hazards	Genetic Risk	chronic Lung Disease
count	1000.000000	1000.000000	1000.0000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	37.174000	1.402000	3.8400	4.563000	5.165000	4.840000	4.580000	4.380000
std	12.005493	0.490547	2.0304	2.620477	1.980833	2.107805	2.126999	1.848518
min	14.000000	1.000000	1.0000	1.000000	1.000000	1.000000	1.000000	1.000000
25%	27.750000	1.000000	2.0000	2.000000	4.000000	3.000000	2.000000	3.000000
50%	36.000000	1.000000	3.0000	5.000000	6.000000	5.000000	5.000000	4.000000
75%	45.000000	2.000000	6.0000	7.000000	7.000000	7.000000	7.000000	6.000000
max	73.000000	2.000000	8.0000	8.000000	8.000000	8.000000	7.000000	7.000000

In [18]:

```
df2.iloc[:, 12:].describe()
```

Out[18]:

	Chest Pain	Coughing of Blood	Fatigue	Weight Loss	Shortness of Breath	Wheezing	Swallowing Difficulty	Clubbing of Finger Nails
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	4.438000	4.859000	3.856000	3.855000	4.240000	3.777000	3.746000	3.923000
std	2.280209	2.427965	2.244616	2.206546	2.285087	2.041921	2.270383	2.388048
min	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000
25%	2.000000	3.000000	2.000000	2.000000	2.000000	2.000000	2.000000	2.000000
50%	4.000000	4.000000	3.000000	3.000000	4.000000	4.000000	4.000000	4.000000
75%	7.000000	7.000000	5.000000	6.000000	6.000000	5.000000	5.000000	5.000000
max	9.000000	9.000000	9.000000	8.000000	9.000000	8.000000	8.000000	9.000000

In [19]:

```
df2.shape
```

Out[19]:

(1000, 24)

Preparing the second dataset for visualization

In [20]:

```
sns.set(style="whitegrid")
plt.figure(figsize=(10, 6))
df2['Gender'] = df2['Gender'].map({1: 'Male', 2: 'Female'})
df2['Level'] = df2['Level'].astype("category")
```

<Figure size 1000x600 with 0 Axes>


Plotting the age distribution of the patients

In [21]:

```
plt.figure(figsize=(12, 9))
sns.histplot(df2['Age'], bins=50, kde=True, color='skyblue')
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()

mean_age = df2['Age'].mean()
median_age = df2['Age'].median()
mode_age = df2['Age'].mode()[0]

print(f"Mean Age: {mean_age}")
print(f"Median Age: {median_age}")
print(f"Mode Age: {mode_age}")
```



Mean Age: 37.174

Median Age: 36.0

Mode Age: 35


### Plotting the gender distributions of the patients

In [22]:

```
gender_counts = df2["Gender"].value_counts()

colors = ["#dcbeaf", "#e7e0ee"]
explode = (0.05, 0)
```

```
plt.figure(figsize=(6, 6))
plt.pie(
    gender_counts,
    labels=gender_counts.index,
    autopct="%1.1f%%",
    startangle=90,
    colors=colors,
    explode=explode,
    shadow=True
)
plt.title("Gender Distribution")
plt.axis("equal")
plt.show()
```



## Plotting the risks of Lung Cancer of the patients

In [23]:

```
sns.countplot(data=df2, x='Level', palette='Set2')
plt.title('Lung Cancer Risk Levels')
plt.show();
```


<ipython-input-23-80d4d314cecc>:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(data=df2, x='Level', palette='Set2')
```

## Lung Cancer Risk Levels





Age vs Lung Cancer Boxplot


In [24]:

```
plt.figure(figsize=(10, 6))
sns.boxplot(data=df2, x='Level', y='Age', palette='coolwarm')
plt.title('Age vs Lung Cancer Risk Level')
plt.show()
```

<ipython-input-24-db9e4c8384cb>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df2, x='Level', y='Age', palette='coolwarm')
```



High

Low  
Level

Medium

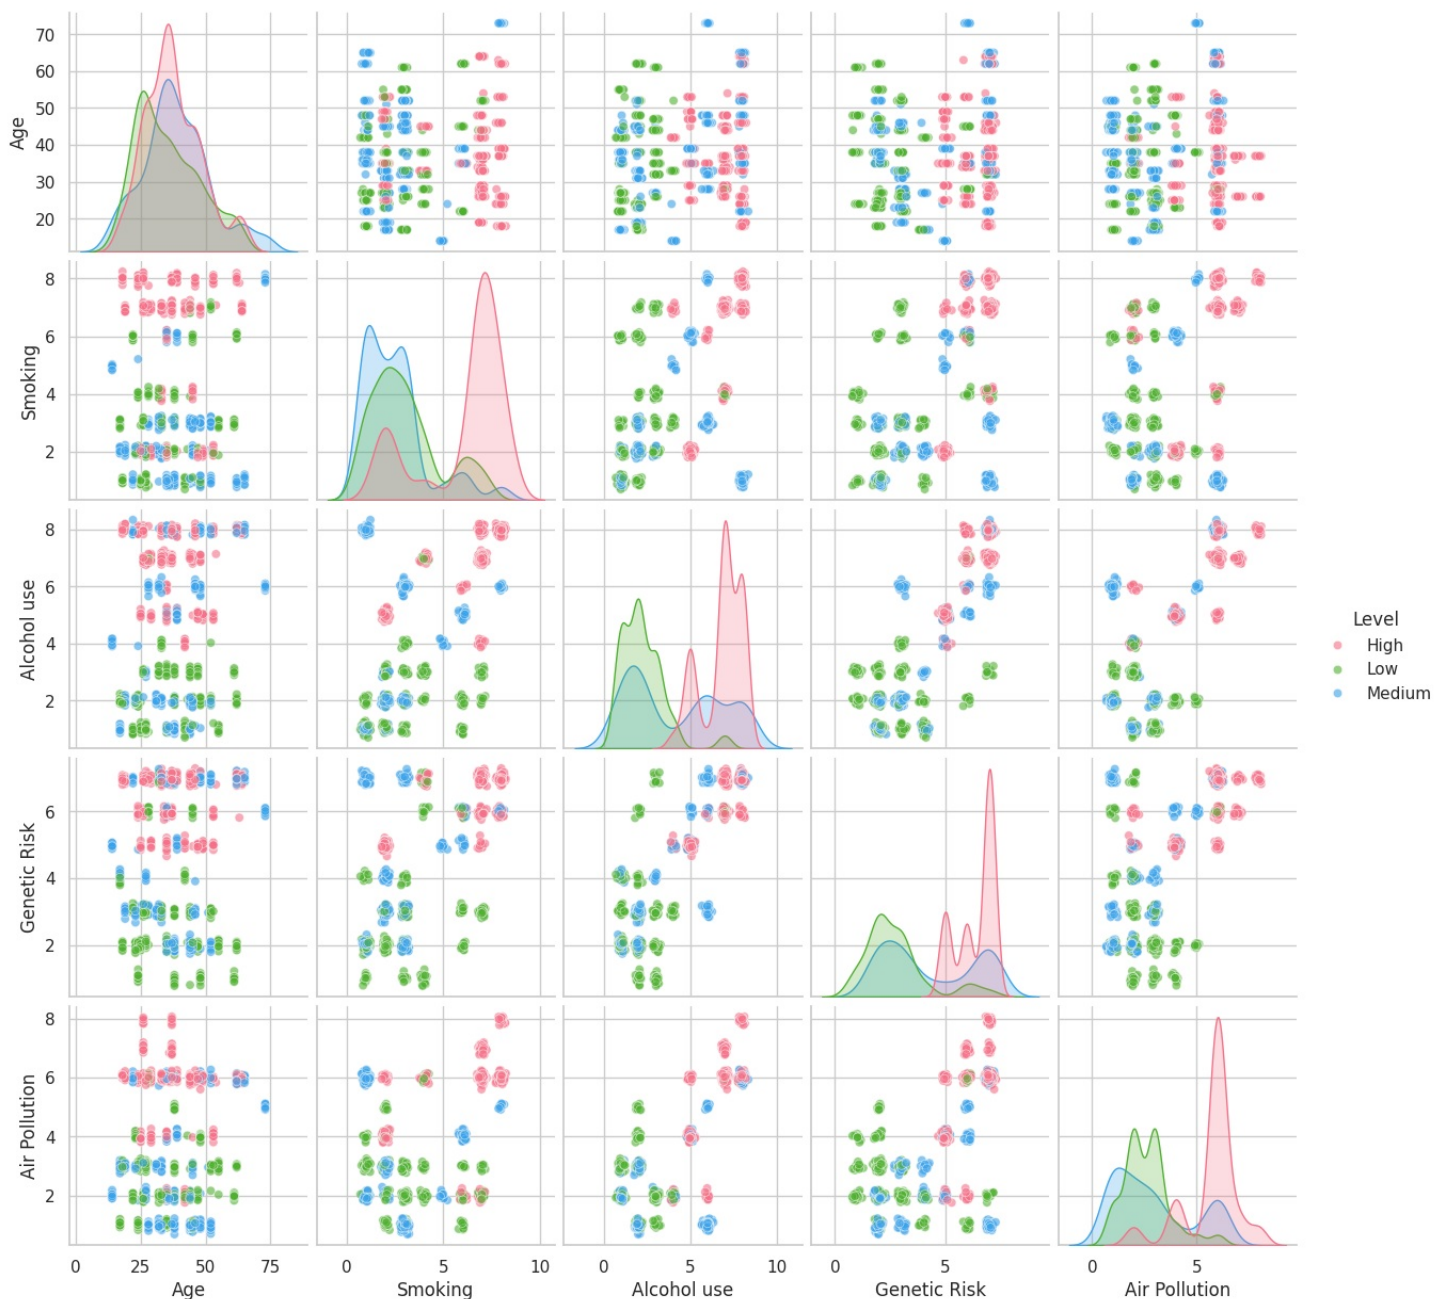
## Pairplot of the selected columns

In [25]:

```
selected = ['Age', 'Smoking', 'Alcohol use', 'Genetic Risk', 'Air Pollution', 'Level']
df_jittered = df2.copy()
jitter_cols = ['Smoking', 'Alcohol use', 'Genetic Risk', 'Air Pollution']

for col in jitter_cols:
    df_jittered[col] = df_jittered[col] + np.random.normal(0, 0.1, size=len(df2))

sns.pairplot(df_jittered[selected], hue='Level', palette='husl', plot_kws={'alpha': 0.6});
```




## Correlation between the columns

In [26]:

```
categorical_cols = df2.drop(columns=['Age', 'Gender', 'Level']).columns
plt.figure(figsize=(12, 10))
```



```
corr_matrix = df2[categorical_cols].corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", square=True)
plt.title("Correlation Between Categorical Risk Factors")
plt.show()
```



## Hypothesis Tests

This section will be about some hypothesis on the causes of lung cancer. The two datasets will be compared using hypothesis tests on the same hypotheses.

### 1. Hypothesis: Smoking regularly affects the risk of lung cancer.

- **Null Hypothesis ( $H_0$ ):** Smoking and lung cancer are independent
- **Alternative Hypothesis ( $H_1$ ):** Smoking and lung cancer are associated

### Dataset 1

In [27]:

```
df["SMOKING"] = df["SMOKING"].map({"Yes": 1, "No": 0})
df["LUNG_CANCER"] = df["LUNG_CANCER"].map({"YES": 1, "NO": 0})
```



```
contingency_table = pd.crosstab(df["SMOKING"], df["LUNG_CANCER"])

# --- Run chi-square test ---
chi2, p, dof, expected = chi2_contingency(contingency_table)

# --- Print results ---
print("Contingency Table:\n", contingency_table)
print("\nChi-square Statistic:", chi2)
print("Degrees of Freedom:", dof)
print("P-value:", p)

# --- Interpret result ---
alpha = 0.05
if p < alpha:
    print("Reject the null hypothesis: Smoking is associated with lung cancer.")
else:
    print("Fail to reject the null hypothesis: No significant association between smoking and lung cancer.")
```

Contingency Table:

LUNG_CANCER	0	1
SMOKING		
0	765	762
1	717	756

Chi-square Statistic: 0.5510111488430194

Degrees of Freedom: 1

P-value: 0.45790483338496013

Fail to reject the null hypothesis: No significant association between smoking and lung cancer.

## Dataset 2

In [28]:

```
df2["Level"] = df2["Level"].map({"Low": 0, "Medium": 1, "High": 2})
contingency_table = pd.crosstab(df2['Smoking'], df2['Level'])

# Perform the Chi-Square test
chi2, p_value, dof, expected = chi2_contingency(contingency_table)

# Output results
print("Contingency Table:\n", contingency_table)
print("Chi-Square Statistic:", chi2)
print("Degrees of Freedom:", dof)
print("P-Value:", p_value)

# Interpret the result
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: Smoking is associated with lung cancer.")
else:
    print("Fail to reject the null hypothesis: No significant association between smoking and lung cancer.")
```

Contingency Table:

Level	2	0	1
Smoking			
1	0	61	120
2	70	81	71
3	0	71	101
4	19	40	0
5	0	0	10
6	10	0	0

```

10    30    20
7      187   20    0
8       79    0   10
Chi-Square Statistic: 684.4965209399204
Degrees of Freedom: 14
P-Value: 5.246169667913276e-137
Reject the null hypothesis: Smoking is associated with lung cancer.

```

**According to the first dataset, smoking regularly does not have a significant relation with lung cancer. However, the test for the second dataset tells us the otherwise.**

### 1. Hypothesis: Having chronic lung diseases affect the risk of lung cancer.

- **Null Hypothesis ( $H_0$ ):** Chronic lung diseases and lung cancer are independent
- **Alternative Hypothesis ( $H_1$ ):** Chronic lung diseases and lung cancer are associated

### Dataset 1

In [29]:

```

df["CHRONIC_DISEASE"] = df["CHRONIC_DISEASE"].map({"Yes": 1, "No": 0})

contingency_table = pd.crosstab(df["CHRONIC_DISEASE"], df["LUNG_CANCER"])

# --- Run chi-square test ---
chi2, p, dof, expected = chi2_contingency(contingency_table)

# --- Print results ---
print("Contingency Table:\n", contingency_table)
print("\nChi-square Statistic:", chi2)
print("Degrees of Freedom:", dof)
print("P-value:", p)

# --- Interpret result ---
alpha = 0.05
if p < alpha:
    print("Reject the null hypothesis: Chronic diseases are associated with lung cancer.")
else:
    print("Fail to reject the null hypothesis: No significant association between chronic diseases and lung cancer.")

```

```

Contingency Table:
LUNG_CANCER      0      1
CHRONIC_DISEASE
0              719   752
1              763   766

```

```

Chi-square Statistic: 0.27462898733847857
Degrees of Freedom: 1
P-value: 0.6002433756617297
Fail to reject the null hypothesis: No significant association between chronic diseases and lung cancer.

```

### Dataset 2

In [30]:

```

contingency_table = pd.crosstab(df2['chronic Lung Disease'], df2['Level'])

# Perform the Chi-Square test

```

```
chi2, p_value, dof, expected = chi2_contingency(contingency_table)
```

```
# Output results
```

```
print("Contingency Table:\n", contingency_table)
```

```
print("Chi-Square Statistic:", chi2)
```

```
print("Degrees of Freedom:", dof)
```

```
print("P-Value:", p_value)
```

```
# Interpret the result
```

```
alpha = 0.05
```

```
if p_value < alpha:
```

```
    print("Reject the null hypothesis: Chronic diseases are associated with lung cancer.")
```

```
else:
```

```
    print("Fail to reject the null hypothesis: No significant association between chronic diseases and lung cancer.")
```

Contingency Table:

Level	2	0	1
chronic Lung Disease			
1	0	50	0
2	0	82	91
3	0	81	60
4	70	20	51
5	10	40	30
6	198	10	100
7	87	20	0

Chi-Square Statistic: 585.1295399409582

Degrees of Freedom: 12

P-Value: 1.5853456732363079e-117

Reject the null hypothesis: Chronic diseases are associated with lung cancer.

**According to the first dataset, chronic lung diseases regularly do not have a significant relation with lung cancer. However, the test for the second dataset tells us the otherwise.**

### 1. Hypothesis: Gender has a relation with lung cancer.

- **Null Hypothesis ( $H_0$ ):** Gender and lung cancer is independent
- **Alternative Hypothesis ( $H_1$ ):** Gender and lung cancer is associated

### Dataset 1

In [31]:

```
df["GENDER"] = df["GENDER"].map({"Male": 1, "Female": 2})
```

```
contingency_table = pd.crosstab(df["GENDER"], df["LUNG_CANCER"])
```

```
# --- Run chi-square test ---
```

```
chi2, p, dof, expected = chi2_contingency(contingency_table)
```

```
# --- Print results ---
```

```
print("Contingency Table:\n", contingency_table)
```

```
print("\nChi-square Statistic:", chi2)
```

```
print("Degrees of Freedom:", dof)
```

```
print("P-value:", p)
```

```
# --- Interpret result ---
```

```
alpha = 0.05
```

```
if p < alpha:
```

```
    print("Reject the null hypothesis: Gender is associated with lung cancer.")
```

```
else:
    print("Fail to reject the null hypothesis: No significant association between
gender and lung cancer.")
```

Contingency Table:

	LUNG_CANCER	0	1
GENDER			
1		760	754
2		722	764

Chi-square Statistic: 0.7158403950228966

Degrees of Freedom: 1

P-value: 0.3975117338177119

Fail to reject the null hypothesis: No significant association between gender and lung cancer.

## Dataset 2

In [32]:

```
contingency_table = pd.crosstab(df2['Gender'], df2['Level'])

# Perform the Chi-Square test
chi2, p_value, dof, expected = chi2_contingency(contingency_table)

# Output results
print("Contingency Table:\n", contingency_table)
print("Chi-Square Statistic:", chi2)
print("Degrees of Freedom:", dof)
print("P-Value:", p_value)

# Interpret the result
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: Gender is associated with lung cancer.")
else:
    print("Fail to reject the null hypothesis: No significant association between
gender and lung cancer.")
```

Contingency Table:

	Level	2	0	1
Gender				
Female		113	154	135
Male		252	149	197

Chi-Square Statistic: 27.22494142912069

Degrees of Freedom: 2

P-Value: 1.2251212802771069e-06

Reject the null hypothesis: Gender is associated with lung cancer.

**According to the first dataset, gender regularly do not have a significant relation with lung cancer. However, the test for the second dataset tells us the otherwise.**

Even though all of the test results were different for both datasets, the reason may only be the sizes of them. If both of the datasets had 10k+ patients, the result might have been similar.

## Creating machine learning models using sklearn

In [33]:

```
from sklearn.model_selection import train_test_split
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix,
roc_curve, auc
from sklearn.preprocessing import StandardScaler, LabelEncoder

# Set seaborn style
sns.set(style="whitegrid")

```

In [34]:

```
df.head()
```

Out[34]:

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CHRONIC_DISEASE	FATIGUE	ALLI
0	1	65	0	No	No	Yes	1	No	
1	2	55	0	Yes	Yes	No	0	Yes	
2	2	78	1	Yes	No	No	0	Yes	
3	1	60	1	No	No	No	1	No	
4	2	80	0	No	Yes	No	0	Yes	

In [35]:

```
df2.head()
```

Out[35]:

	Age	Gender	Air Pollution	Alcohol use	Dust Allergy	OccuPational Hazards	Genetic Risk	chronic Lung Disease	Balanced Diet	Obesity	...	Fatigue	Weigh Lo
0	33	Male	2	4	5	4	3	2	2	4	...	3	
1	17	Male	3	1	5	3	4	2	2	2	...	1	
2	35	Male	4	5	6	5	5	4	6	7	...	8	
3	37	Male	7	7	7	7	6	7	7	7	...	4	
4	46	Male	6	8	7	7	7	6	7	7	...	3	

5 rows x 24 columns

In [36]:

```

# Replace categorical 'Yes'/'No' and '1'/'2' with 1/0
df = df.replace({"Yes": 1, "No": 0, "2": 1, "1": 0})

# Split features (X) and target (y)
X = df.drop("LUNG_CANCER", axis=1)
y = df["LUNG_CANCER"]

# Check the distribution of the target variable
print(f"Target distribution:\n{y.value_counts()}")

```

Target distribution:

LUNG\_CANCER

1 1518

```
1 1510
0 1482
Name: count, dtype: int64
```

```
<ipython-input-36-e98bae2b7c38>:2: FutureWarning: Downcasting behavior in `replace`
is deprecated and will be removed in a future version. To retain the old behavior,
explicitly call `result.infer_objects(copy=False)`. To opt-in to the future
behavior, set `pd.set_option('future.no_silent_downcasting', True)`
df = df.replace({"Yes": 1, "No": 0, "2": 1, "1": 0})
```

In [37]:

```
df.head()
```

Out[37]:

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CHRONIC_DISEASE	FATIGUE	ALLI
0	1	65	0	0	0	1	1	0	
1	2	55	0	1	1	0	0	1	
2	2	78	1	1	0	0	0	1	
3	1	60	1	0	0	0	1	0	
4	2	80	0	0	1	0	0	1	

In [38]:

```
# Split the data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Check the shape of train/test datasets
print(f"Training set: {X_train.shape}")
print(f"Testing set: {X_test.shape}")
```

```
Training set: (2400, 15)
Testing set: (600, 15)
```

This section of the code focuses on training and evaluating several common machine learning classification models using the prepared data. The goal is to see which model performs best at predicting the target variable, which in this case is likely related to lung cancer status.

In [39]:

```
%%time
# Define the classification models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC(probability=True),
    "k-NN": KNeighborsClassifier()
}

# Store results
results = {}
y_probs = {}

# Train each model and evaluate performance
for name, model in models.items():
    model.fit(X_train, y_train)
```

```

y_pred = model.predict(X_test)
acc = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

results[name] = {
    "accuracy": acc,
    "classification_report": report,
    "conf_matrix": confusion_matrix(y_test, y_pred)
}

if hasattr(model, "predict_proba"):
    y_probs[name] = model.predict_proba(X_test)[:, 1]

print(f"\n{name}")
print(f"Accuracy: {acc:.2f}")
print("Classification Report:\n", report)

```

#### Logistic Regression

Accuracy: 0.52

Classification Report:

	precision	recall	f1-score	support
0	0.52	0.45	0.48	296
1	0.53	0.59	0.56	304
accuracy			0.52	600
macro avg	0.52	0.52	0.52	600
weighted avg	0.52	0.52	0.52	600

#### Random Forest

Accuracy: 0.55

Classification Report:

	precision	recall	f1-score	support
0	0.54	0.54	0.54	296
1	0.55	0.56	0.56	304
accuracy			0.55	600
macro avg	0.55	0.55	0.55	600
weighted avg	0.55	0.55	0.55	600

#### SVM

Accuracy: 0.52

Classification Report:

	precision	recall	f1-score	support
0	0.52	0.44	0.48	296
1	0.52	0.60	0.56	304
accuracy			0.52	600
macro avg	0.52	0.52	0.52	600
weighted avg	0.52	0.52	0.52	600

#### k-NN

Accuracy: 0.55

Classification Report:

	precision	recall	f1-score	support
0	0.54	0.56	0.55	296
1	0.55	0.54	0.55	304
accuracy			0.55	600

macro avg	0.55	0.55	0.55	600
weighted avg	0.55	0.55	0.55	600

CPU times: user 2.34 s, sys: 79 ms, total: 2.41 s

Wall time: 2.29 s


In [40]:

```
# Plotting model accuracy comparison
accuracies = [results[m]["accuracy"] for m in models]
plt.figure(figsize=(8, 5))
sns.barplot(x=list(models.keys()), y=accuracies, palette="viridis")
plt.title("Model Accuracy Comparison")
plt.ylabel("Accuracy")
plt.ylim(0, 1)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

<ipython-input-40-5f2253f25700>:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=list(models.keys()), y=accuracies, palette="viridis")
```



In [41]:

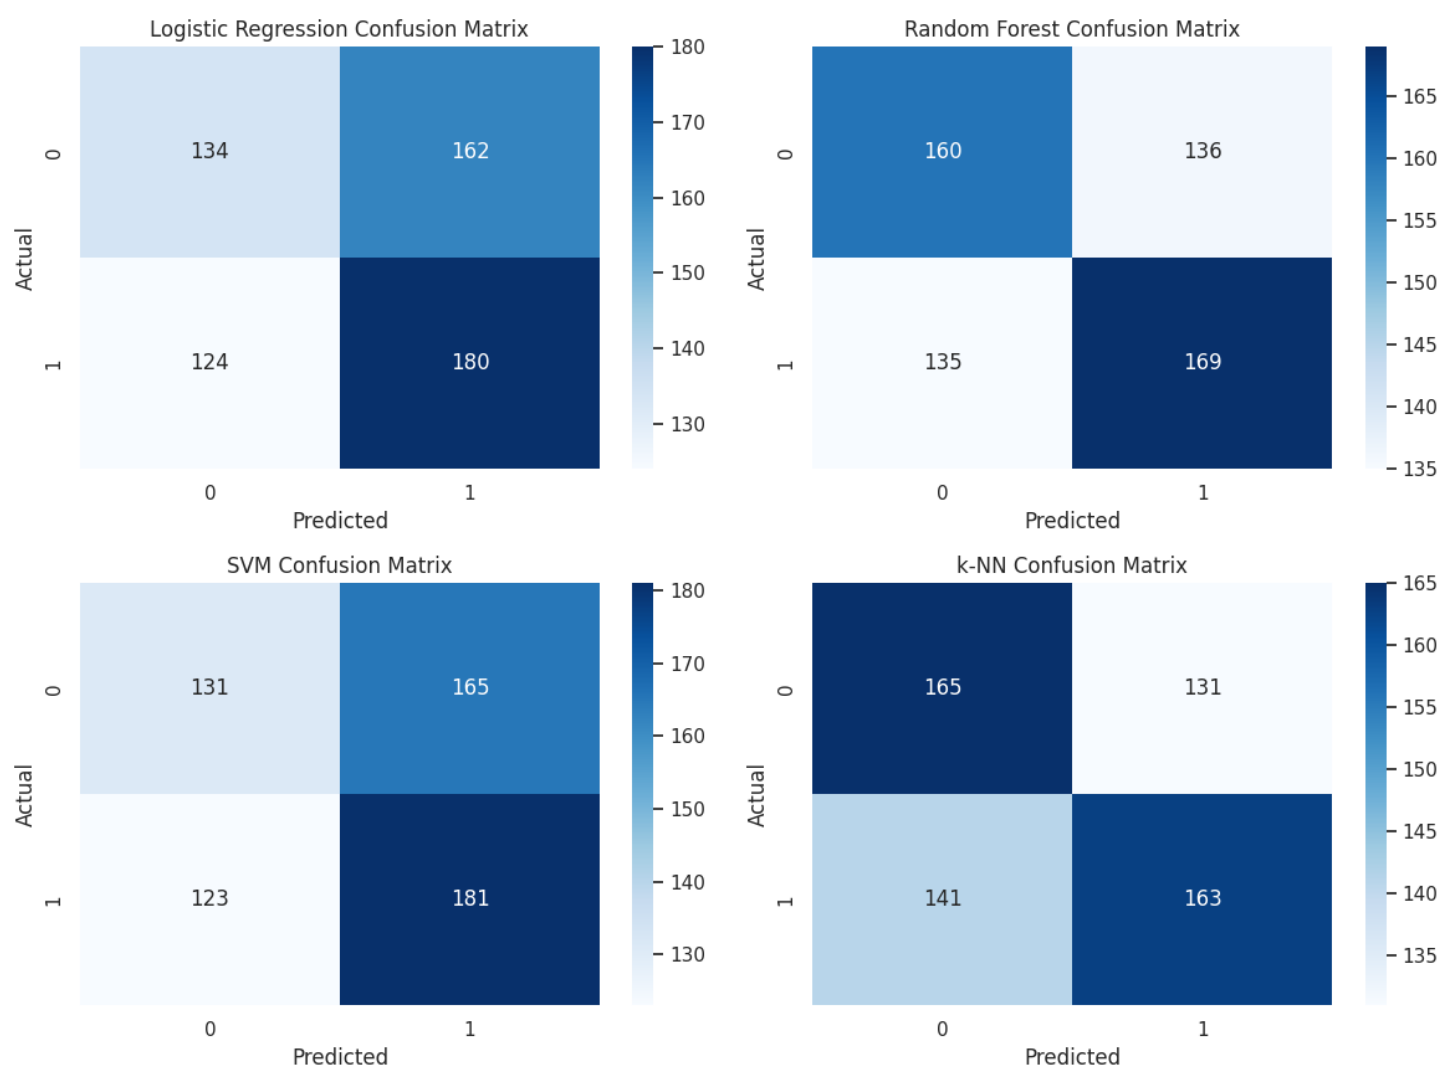
```
# Plot confusion matrices for each model
fig, axes = plt.subplots(2, 2, figsize=(12, 10))
axes = axes.flatten()

for idx, name in enumerate(models):
    sns.heatmap(results[name]["conf_matrix"], annot=True, fmt="d", cmap="Blues", ax=
axes[idx])
    axes[idx].set_title(f"{name} Confusion Matrix")
    axes[idx].set_xlabel("Predicted")
    axes[idx].set_ylabel("Actual")
```



```
plt.suptitle("Confusion Matrices", fontsize=16)
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```

Confusion Matrices



In [42]:


```
# Plot ROC curves for each model
plt.figure(figsize=(8, 6))
for name, probs in y_probs.items():
    fpr, tpr, _ = roc_curve(y_test, probs)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f"{name} (AUC = {roc_auc:.2f})")


# Random line for comparison
plt.plot([0, 1], [0, 1], "k--", label="Random")
plt.title("ROC Curves")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid()
plt.tight_layout()
plt.show()
```

ROC Curves

1.0

— Logistic Regression (AUC = 0.53)  
— Random Forest (AUC = 0.54)  
— SVM (AUC = 0.45)  
— k-NN (AUC = 0.55)





Overall results show that different types of training models of the first dataset have between 0.50-0.55 accuracy, highest being 0.55 using k-NN as a the base model. Since the first dataset was relatively simple (as we can see from the correlation matrix), this output was expected.

In [43]:

```
df2['Gender'] = df2['Gender'].map({'Male': 1, 'Female': 2})
```

```
# Split features (X) and target (y)
```

```
X = df2.drop("Level", axis=1)
```

```
y = df2["Level"]
```

```
# Check the distribution of the target variable
```

```
print(f"Target distribution:\n{y.value_counts()}")
```

```
Target distribution:
```

```
Level
```

```
2      365
```

```
1      332
```

```
0      303
```

```
Name: count, dtype: int64
```

In [44]:

```
# Split the data into training and testing sets (80% training, 20% testing)
```

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y, test_size=0.2, random_state=42, stratify=y
```

```
)
```

```
# Check the shape of train/test datasets
```

```
print(f"Training set: {X_train.shape}")
```

```
print(f"Testing set: {X_test.shape}")
```

```
Training set: (800, 23)
```

```
Testing set: (200, 23)
```

This section of the code trains and evaluates several different machine learning models on the second data. The goal is to see how well each model can predict the target variable.

data. The goal is to see how well each model can predict the target variable.

In [45]:

```
%%time
# Define the classification models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC(probability=True),
    "k-NN": KNeighborsClassifier()
}

# Store results
results = {}
y_probs = {}

# Train each model and evaluate performance
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    report = classification_report(y_test, y_pred)

    results[name] = {
        "accuracy": acc,
        "classification_report": report,
        "conf_matrix": confusion_matrix(y_test, y_pred)
    }

    if hasattr(model, "predict_proba"):
        y_probs[name] = model.predict_proba(X_test)[:, 1]

    print(f"\n{name}")
    print(f"Accuracy: {acc:.2f}")
    print("Classification Report:\n", report)
```

Logistic Regression

Accuracy: 1.00

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	61
1	1.00	1.00	1.00	66
2	1.00	1.00	1.00	73
accuracy			1.00	200
macro avg	1.00	1.00	1.00	200
weighted avg	1.00	1.00	1.00	200

Random Forest

Accuracy: 1.00

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	61
1	1.00	1.00	1.00	66
2	1.00	1.00	1.00	73
accuracy			1.00	200
macro avg	1.00	1.00	1.00	200
weighted avg	1.00	1.00	1.00	200

SVM

Accuracy: 0.98

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.97	0.98	61
1	0.97	0.98	0.98	66
2	0.99	1.00	0.99	73
accuracy			0.98	200
macro avg	0.99	0.98	0.98	200
weighted avg	0.99	0.98	0.98	200

k-NN

Accuracy: 0.99

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.97	0.98	61
1	0.97	1.00	0.99	66
2	1.00	1.00	1.00	73
accuracy			0.99	200
macro avg	0.99	0.99	0.99	200
weighted avg	0.99	0.99	0.99	200

CPU times: user 2.76 s, sys: 8.33 ms, total: 2.77 s

Wall time: 2.43 s


In [46]:

```
# Plotting model accuracy comparison
accuracies = [results[m]["accuracy"] for m in models]
plt.figure(figsize=(8, 5))
sns.barplot(x=list(models.keys()), y=accuracies, palette="viridis")
plt.title("Model Accuracy Comparison")
plt.ylabel("Accuracy")
plt.ylim(0, 1)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

<ipython-input-46-5f2253f25700>:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=list(models.keys()), y=accuracies, palette="viridis")
```





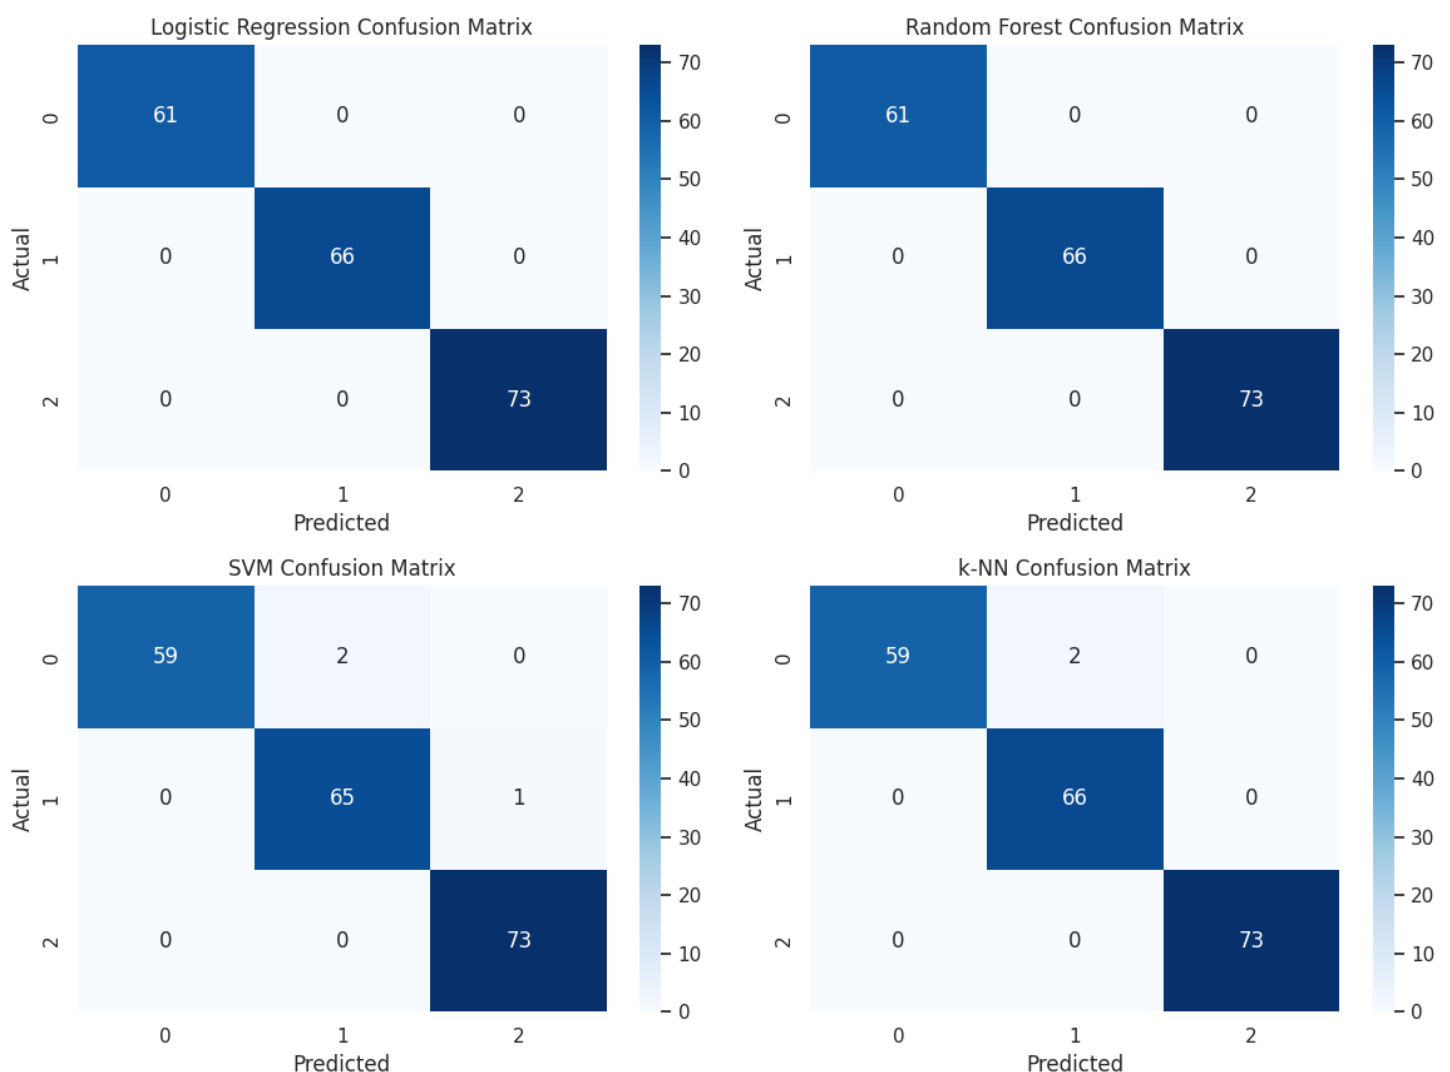
In [47]:

```
# Plot confusion matrices for each model
fig, axes = plt.subplots(2, 2, figsize=(12, 10))
axes = axes.flatten()

for idx, name in enumerate(models):
    sns.heatmap(results[name]["conf_matrix"], annot=True, fmt="d", cmap="Blues", ax=
axes[idx])
    axes[idx].set_title(f"{name} Confusion Matrix")
    axes[idx].set_xlabel("Predicted")
    axes[idx].set_ylabel("Actual")

plt.suptitle("Confusion Matrices", fontsize=16)
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```

Confusion Matrices



In [48]:

```

from sklearn.preprocessing import label_binarize

# Binarize the output for multiclass ROC
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))
n_classes = y_test_bin.shape[1]

# Prepare plot
plt.figure(figsize=(8, 6))

# Plot ROC for each model
for name, model in models.items():
    if hasattr(model, "predict_proba"):
        y_score = model.predict_proba(X_test)
    elif hasattr(model, "decision_function"):
        y_score = model.decision_function(X_test)
        # If 1D, reshape to 2D
        if y_score.ndim == 1:
            y_score = y_score.reshape(-1, 1)
    else:
        print(f"{name} does not support probability/decision score; skipping ROC.")
        continue

    if n_classes == 1:
        print(f"{name}: Only one class present; skipping.")
        continue

    # Compute ROC curve and AUC for each class
    fpr = dict()
    tpr = dict()
    roc_auc = dict()
    for i in range(n_classes):
        fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
        roc_auc[i] = auc(fpr[i], tpr[i])

    # Average macro AUC
    all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))
    mean_tpr = np.zeros_like(all_fpr)
    for i in range(n_classes):
        mean_tpr += np.interp(all_fpr, fpr[i], tpr[i])
    mean_tpr /= n_classes

    roc_auc["macro"] = auc(all_fpr, mean_tpr)

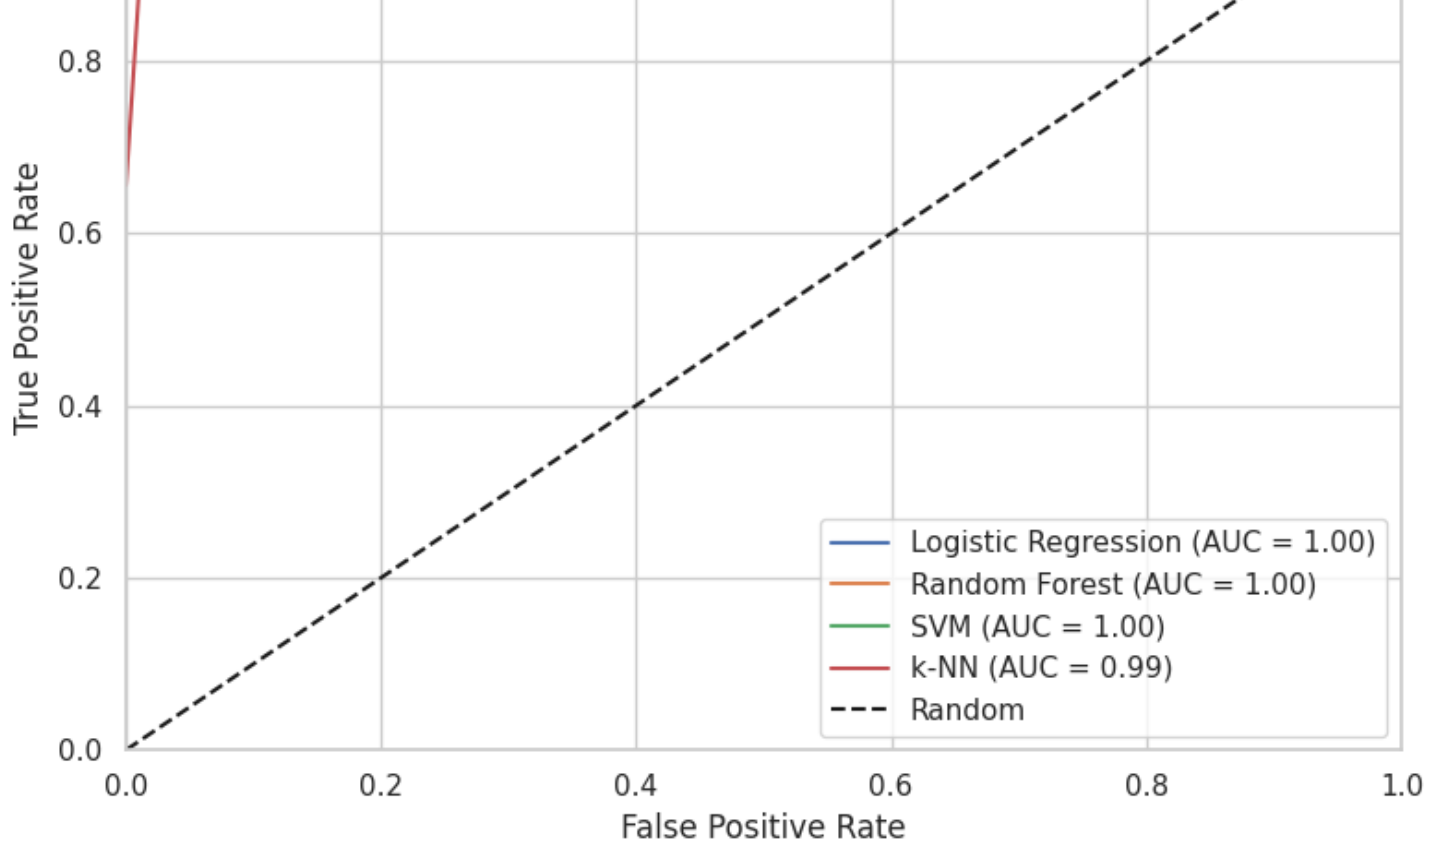
    # Plot macro-average ROC
    plt.plot(all_fpr, mean_tpr, label=f"{name} (AUC = {roc_auc['macro']:.2f})")

# Plot settings
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves for All Models')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

```

ROC Curves for All Models





Overall results show that different types of training models of the second dataset have between 0.98-1.00 accuracy, highest being 1.00 using Logistic Regression and Random Forest as the base model. The result was unexpected and the accuracy level of 1.00 indicates an error most of the times. Since the dataset used can be considered small, a certain column could have had an effect on the output. This could be checked using correlation matrix.

## RESULTS

This concludes the comparison between the datasets. As a result, the second dataset achieved far better solutions from my expectations. The main reason is most likely the lack of data in the dataset. If both datasets had more data to process on, the results would differ as well. Nevertheless, the outputs show that the second dataset is better for lung cancer detection.

In [48]: